

Agent observability & evaluation tools

Pros and cons comparison · compiled July 2026

Recommended tools (overall)

Tool	Pros	Cons
LangSmith	Zero-config tracing for LangChain/LangGraph; great annotation queues; polished UI	Best value only within LangChain ecosystem; closed source; per-seat costs
Langfuse	Most popular OSS option; self-hostable; prompt management + evals built in	Less polished UI; needs ClickHouse/Postgres/Redis to self-host
Braintrust	Evaluation-first; 25+ scorers; evals gate releases via GitHub Actions	Closed source; self-host is enterprise-only; \$249/mo Pro
Arize Phoenix / AX	Strong OpenTelemetry base; drift detection; enterprise-scale	Thinner built-in LLM eval metrics; confusing OSS vs paid split
MLflow	Fully free, no paywalls; built-in LLM judges; covers ML + GenAI	You run all infra yourself; agent features are newer
Galileo	Sub-200ms evaluators, ~\$0.02/M tokens; cheap real-time guardrails	Closed source; small free tier; guardrails-focused
Confident AI	50+ eval metrics on every trace; quality alerts; PM/QA friendly	Smaller vendor; costs hard to forecast; heavy self-marketing
Datadog / New Relic	Sits beside your existing APM; correlate LLM + infra issues	Needs existing subscription; pricey at scale; evals bolted on

Open source options

Tool	License	Pros	Cons
Langfuse	MIT (core)	Full platform: tracing, prompts, evals; 5-min Docker setup; OTel support	SSO/RBAC separately licensed; multi-service self-host stack
Agenta	MIT	Self-hostable; trace annotation; LLM-as-judge evals	Small community; manual, request-level monitoring
TruLens	MIT*	Strong RAG metrics; runs fully local; evaluates agent steps	Evaluation only — needs a separate tracing tool
Promptfoo	MIT*	CLI eval testing in CI; includes red-teaming	Testing harness only, no production tracing
Opik (Comet)	Apache 2.0	Most complete OSS agent platform; testing + auto-optimization	Newer; smaller community
MLflow	Apache 2.0	Everything free, Linux Foundation-governed; AI gateway included	Self-managed infra; no polished managed cloud
Laminar	Apache 2.0	Built for long-running agents; one-command Helm deploy	Younger project; smaller ecosystem
OpenLLMetry	Apache 2.0	OTel-native instrumentation; backend-agnostic	Instrumentation layer only, no UI/platform

Tool	License	Pros	Cons
Arize Phoenix	Elastic 2.0	Solid tracing + evals; OTel-friendly	Source-available, not true OSS; can't resell as a service

* Verify the current license in each repository before committing — licenses occasionally change.

Quick picks: Langfuse for one self-hosted platform · Opik if Apache 2.0 is fine and you want the richest agent tooling · Braintrust or LangSmith if you're happy paying for polish.